

Assignment 06

Neural Networks – Part 1



Assignment 06

Solution

1. Activation Functions

2. Neural networks

Activation Functions



Activation Functions

Task 1.1: Non-Linear vs. Linear

Why do we use non-linear activation functions?

- Stacking multiple linear functions will again create a linear function (collapse).
- Example:

2 Layer Linear Network

$$\begin{aligned}h^{(1)} &= W^{(1)} \cdot x + b^{(1)} \\h^{(2)} &= W^{(2)} \cdot h^{(1)} + b^{(2)}\end{aligned}$$

Equivalent single-layer Network

$$\begin{aligned}h^{(2)} &= W^{(2)} \cdot (W^{(1)} \cdot x + b^{(1)}) + b^{(2)} \\&= (W^{(2)} W^{(1)}) \cdot x + (W^{(2)} b^{(1)}) + b^{(2)} \\&= W' \cdot x + b'\end{aligned}$$

Activation Functions

Task 1.2: Computing Gradients

Gradient for ReLU

$$\begin{aligned}\text{ReLU}(z) &= \max(z, 0) \\ &= \begin{cases} z & z > 0 \\ 0 & z \leq 0 \end{cases}\end{aligned}\quad \frac{\partial \text{ReLU}(z)}{\partial z} = \begin{cases} 1 & z > 0 \\ 0 & z < 0 \end{cases}$$

Undefined for $z = 0$!

Activation Functions

Task 1.2: Computing Gradients

Gradient for Sigmoid

- Start with basic derivation:

$$\sigma(z) = \frac{1}{1+e^{-z}} \quad \frac{\partial \sigma(z)}{\partial z} = \frac{\partial}{\partial z} \left(\frac{1}{1+e^{-z}} \right) = -1 \cdot (1+e^{-z})^{-2} \cdot (-e^{-z}) = \frac{e^{-z}}{(1+e^{-z})^2}$$

- Express terms via $\sigma(z)$:

$$\sigma(z) = \frac{1}{1+e^{-z}}, \quad 1+e^{-z} = \frac{1}{\sigma(z)}, \quad e^{-z} = \frac{1-\sigma(z)}{\sigma(z)}$$

- Then simplify:

$$\frac{\partial \sigma(z)}{\partial z} = \frac{1-\sigma(z)}{\sigma(z)} \cdot (\sigma(z))^2 = (1-\sigma(z)) \cdot \sigma(z)$$

Activation Functions

Task 1.2: Computing Gradients

Gradient for TanH

$$\begin{aligned}\tanh(z) &= \frac{e^{2z}-1}{e^{2z}+1} & \frac{\partial \tanh(z)}{\partial z} &= \frac{\partial}{\partial z} \left(\frac{e^{2z}-1}{e^{2z}+1} \right) && \text{apply quotient rule} \\ && &= \frac{2e^{2z} \cdot (e^{2z}+1) - (e^{2z}-1) \cdot 2e^{2z}}{(e^{2z}+1)^2} && \text{multiply out} \\ && &= \frac{2e^{4z} + 2e^{2z} - 2e^{4z} + 2e^{2z}}{(e^{2z}+1)^2} && \text{subtract } e^{4z} \text{ and expand by } +/-1 \\ && &= \frac{e^{4z} + 2e^{2z} + 1 - e^{4z} + 2e^{2z} - 1}{(e^{2z}+1)^2} && \text{separate} \\ && &= \frac{(e^{4z} + 2e^{2z} + 1) - (e^{4z} - 2e^{2z} + 1)}{(e^{2z}+1)^2} && \text{apply square rule} \\ && &= \frac{(e^{2z}+1)^2 - (e^{2z}-1)^2}{(e^{2z}+1)^2} = 1 - \tanh^2(z)\end{aligned}$$

Activation Functions

Task 1.3: Advantages and Disadvantages

Advantages and disadvantages of the ReLU activation function

- Fixes vanishing gradient problem
- Computationally inexpensive
- Fragile during training, can “die”
 - Can be overcome with Leaky ReLU, Parametric ReLU

Activation Functions

Task 1.3: Advantages and Disadvantages

Advantages and disadvantages of the Sigmoid activation function

- Computationally more expensive than ReLU
- Not zero-centered
- Has a smooth gradient
- Well-suited for binary classification when used on the output layer (cf. Logistic regression)
- Suffers from vanishing gradient problems

Activation Functions

Task 1.3: Advantages and Disadvantages

Advantages and disadvantages of the TanH activation function

- Computationally more expensive than ReLU
- Is zero-centered
- Has a smooth gradient
- Used more frequently compared to the sigmoid as an activation function in hidden layers.
 - Special use in recurrent networks (LSTM, GRU)
- Suffers from vanishing gradient problems

Activation Functions

Task 1.4: Softmax Predictions

Given a single data point x with:

true label	possible labels	unnormalized log probability
$y=A$	$Y=\{A,B,C\}$	$z_A=7, \quad z_B=1, \quad z_C=2.2$

Compute softmax output:

$$P(y=i \mid x) \approx \text{softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$
$$A: \text{softmax}(z_A) = \frac{e^{z_A}}{e^{z_A} + e^{z_B} + e^{z_C}} \approx \frac{1097}{1097 + 2.72 + 9.03} \approx 98.9\%$$
$$B: \text{softmax}(z_B) = \frac{e^{z_B}}{e^{z_A} + e^{z_B} + e^{z_C}} \approx \frac{2.72}{1097 + 2.72 + 9.03} \approx 0.25\%$$
$$C: \text{softmax}(z_C) = \frac{e^{z_C}}{e^{z_A} + e^{z_B} + e^{z_C}} \approx \frac{9.03}{1097 + 2.72 + 9.03} \approx 0.81\%$$

Activation Functions

Task 1.4: Softmax Predictions

Given a single data point x with:

true label	possible labels	unnormalized log probability
$y = A$	$Y = \{A, B, C\}$	$z_A = 7, \quad z_B = 1, \quad z_C = 2.2$

Predicted class:

$$A: \text{softmax}(z_A) \approx 98.9\%$$

$$B: \text{softmax}(z_B) \approx 0.25\%$$

$$C: \text{softmax}(z_C) \approx 0.81\%$$

The predicted class is the class with the highest softmax score
→ predicted class is A

Activation Functions

Task 1.4: Softmax Predictions

Given a single data point x with:

true label	possible labels	unnormalized log probability
$y=A$	$Y=\{A,B,C\}$	$z_A=7, \quad z_B=1, \quad z_C=2.2$

Log-likelihood loss:

$$A: P(y=A \mid x) \approx 98.9\%$$

$$B: P(y=B \mid x) \approx 0.25\%$$

$$C: P(y=C \mid x) \approx 0.81\%$$

$$\begin{aligned} J(w, x) &= -\sum_{i \in Y} 1_{y_i=i} \cdot \log P(y_i=i \mid x) \\ &= -1 \cdot \log P(y_i=A \mid x) \\ &\approx -1 \cdot \log 0.989 \\ &\approx 0.011 \end{aligned}$$

Neural Networks



Neural Networks

Task 2.1: Basis Functions and Neural Networks

What is the difference between using basis functions and neural networks on input features?

- Basis functions are difficult to design by hand
 - Especially non-linear ones
- Neural networks can be thought of as learning basis functions!
 - Automatic, adapt to the task at hand to achieve optimal performance

Neural Networks

Task 2.2: Compute the Number of Weights in a Neural Network

Given:

- 80 examples
- 6 features
- Labels: {Spring, Summer, Fall, Winter}
- 3 hidden layers, 32 units each

Translate to relevant information:

- 6 input units
- 32 hidden units, 3 hidden layers
- 4 output units

Compute weights per Layer:

- 1st layer: $6 \cdot 32 + 32 = 224$
- 2nd layer: $32 \cdot 32 + 32 = 1056$
- 3rd layer: $32 \cdot 32 + 32 = 1056$
- Output layer: $32 \cdot 4 + 4 = 132$

$\#inputs \cdot \#neurons + \#neurons$
 $\#multiplicative + \#biases$

Compute total number of weights:

- Sum up over layers: 2468 weights

Neural Networks

Task 2.3: Manual Optimization

Dataset:

x	y
1	0
2	0.3
10	1

Network:

$$\begin{aligned} h^{(1)} &= \text{ReLU}(w_{1,1}^{(1)} \cdot x) & w_{1,1}^{(1)} &= -0.1 \\ h^{(2)} &= \text{ReLU}(w_{1,1}^{(2)} \cdot h^{(1)}) & w_{1,1}^{(2)} &= 0.1 \end{aligned}$$

Loss function:

$$J(w) = \frac{1}{N} \sum_{(x,y) \in D} (\hat{y}(x) - y)^2$$

Should we increase or decrease $w_{1,1}^{(1)}$?

- All values x in our dataset are positive
- $w_{1,1}^{(1)}$ is negative
- ➔ ReLU for $h^{(1)}$ will always output zero
- ➔ Network will always output zero
- ➔ Must change $w_{1,1}^{(1)}$ to be positive, e.g.:
 - Current loss is approx. 0.36
 - Set $w_{1,1}^{(1)} = 1$ and loss is approx. 0.02

Neural Networks

Task 2.3: Manual Optimization

Dataset:

x	y
1	0
2	0.3
10	1

Network:

$$\begin{aligned} h^{(1)} &= \text{ReLU}(w_{1,1}^{(1)} \cdot x) & w_{1,1}^{(1)} &= -0.1 \\ h^{(2)} &= \text{ReLU}(w_{1,1}^{(2)} \cdot h^{(1)}) & w_{1,1}^{(2)} &= 0.1 \end{aligned}$$

Loss function:

$$J(w) = \frac{1}{N} \sum_{(x,y) \in D} (\hat{y}(x) - y)^2$$

Bonus: Compute gradient for $w_{1,1}^{(1)}$

- Expand

$$h^{(1)} = \text{ReLU}(a^{(1)}), \quad a^{(1)} = w_{1,1}^{(1)} \cdot x$$

- Apply chain-rule

$$\begin{aligned} \frac{\partial J(w)}{\partial w_{1,1}^{(1)}} &= \frac{\partial J(w)}{\partial h^{(2)}} \cdot \frac{\partial h^{(2)}}{\partial h^{(1)}} \cdot \frac{\partial h^{(1)}}{\partial w_{1,1}^{(1)}} \\ &= \frac{\partial J(w)}{\partial h^{(2)}} \cdot \frac{\partial h^{(2)}}{\partial h^{(1)}} \cdot \frac{\partial h^{(1)}}{\partial a^{(1)}} \cdot \frac{\partial a^{(1)}}{\partial w_{1,1}^{(1)}} \end{aligned}$$

- Note: $a^{(1)}$ will always be negative! Therefore

$$\frac{\partial h^{(1)}}{\partial a^{(1)}} = 0 \quad \text{and} \quad \frac{\partial J(w)}{\partial w_{1,1}^{(1)}} = 0$$

Neural Networks

Task 2.3: Manual Optimization

Dataset:

x	y
1	0
2	0.3
10	1

Network:

$$\begin{aligned} h^{(1)} &= \text{ReLU}(w_{1,1}^{(1)} \cdot x) & w_{1,1}^{(1)} &= -0.1 \\ h^{(2)} &= \text{ReLU}(w_{1,1}^{(2)} \cdot h^{(1)}) & w_{1,1}^{(2)} &= 0.1 \end{aligned}$$

Loss function:

$$J(w) = \frac{1}{N} \sum_{(x,y) \in D} (\hat{y}(x) - y)^2$$

Bonus: Compute gradient for $w_{1,1}^{(1)}$

- Gradient is zero for all samples in our dataset

$$\frac{\partial J(w)}{\partial w_{1,1}^{(1)}} = 0$$

- This means:
 - Any gradient based optimization procedure will not change this weight (for our dataset)!
 - „Dead“ ReLU, network is not learning.
- Note: If we set $w_{1,1}^{(2)} < 0$, this would even be independent of our dataset!

Neural Networks

Task 2.4: Setting up PyTorch

In short:

- Set up a python environment (or use your standard python installation)
 - E.g. [conda](#)
- Then run `pip install torch`

Any problems?

Neural Networks

Task 2.5: Implementing a Network (a)

```
26 def create_2unit_net() -> Module:
27     """
28     ... Create a two-layer MLP (1 hidden layer, 1 output layer) with 2 hidden units
29     ... as described in the exercise.
30
31     ... Returns:
32     ... .. 2-layer MLP module with 2 hidden units.
33     ... """
34     ... # START TODO #####
35     ... # Define the model here
36     ... linear_units = 2
37     ... model = Sequential(
38     ... .. Linear(2, linear_units),
39     ... .. ReLU(),
40     ... .. Linear(linear_units, 2),
41     ... )
42     ... # END TODO #####
43     ... params = model.state_dict()
44
45     ... # Now we assign the model weights since we still did not learn backpropagation
46     ... params['0.weight'] = torch.Tensor(np.array([[3.21, 3.21], [-2.34, -2.34]]))
47     ... params['0.bias'] = torch.Tensor(np.array([-3.21, 2.34]))
48     ... params['2.weight'] = torch.Tensor(np.array([[3.19, 4.64], [-2.68, -3.44]]))
49     ... params['2.bias'] = torch.Tensor(np.array([-4.08, 4.42]))
50     ... model.load_state_dict(params)
51
52     ... return model
```

```
$ python run_2unit_model.py
```

Raw prediction logits:

```
[[ 6.7775993 -3.6295996]
 [-4.08         4.42        ]
 [-4.08         4.42        ]
 [ 6.1599007 -4.1828003]]
```

Prediction after softmax:

```
[[9.9996984e-01 3.0213278e-05]
 [2.0342699e-04 9.9979657e-01]
 [2.0342699e-04 9.9979657e-01]
 [9.9996781e-01 3.2226122e-05]]
```

True labels, one-hot encoded:

```
[[1. 0.]
 [0. 1.]
 [0. 1.]
 [1. 0.]]
```

Loss: 0.00011732114

Neural Networks

Task 2.5: Implementing a Network (b)

```
55 def create_3unit_net() -> Module:
56     """
57     ... Create a two-layer MLP (1 hidden layer, 1 output layer) with 3 hidden units
58     ... as described in the exercise.
59
60     ... Returns:
61     ... .. 2-layer MLP module with 3 hidden units.
62     ... """
63     ... # START TODO #####
64     ... # Define the model here
65     ... linear_units = 3
66     ... model = Sequential(Linear(2, linear_units), ReLU(), Linear(linear_units, 2))
67     ... # END TODO #####
68
69     ... params = model.state_dict()
70     ... # START TODO #####
71     ... # change the model weights
72
73     ... params['0.weight'] = torch.Tensor(np.array([[3.21, 3.21], [-2.34, -2.34], [0, 0]]))
74     ... params['0.bias'] = torch.Tensor(np.array([-3.21, 2.34, 0]))
75     ... params['2.weight'] = torch.Tensor(np.array([[3.19, 4.64, 0], [-2.68, -3.44, 0]]))
76     ... params['2.bias'] = torch.Tensor(np.array([-4.08, 4.42]))
77
78     ... # END TODO #####
79     ... model.load_state_dict(params)
80     ...
81     ... return model
```

```
$ python run_3unit_model.py
```

Raw prediction logits:

```
[[ 6.7775993 -3.6295996]
 [-4.08         4.42        ]
 [-4.08         4.42        ]
 [ 6.1599007 -4.1828003]]
```

Prediction after softmax:

```
[[9.9996984e-01 3.0213278e-05]
 [2.0342699e-04 9.9979657e-01]
 [2.0342699e-04 9.9979657e-01]
 [9.9996781e-01 3.2226122e-05]]
```

True labels, one-hot encoded:

```
[[1. 0.]
 [0. 1.]
 [0. 1.]
 [1. 0.]]
```

Loss: 0.00011732114